Aurora API Gateway - Product Specification

Aurora is a managed API gateway product used by customers to expose, secure, and monitor internal services. This document summarizes the current specification as of the latest release.

Rate Limiting

The gateway enforces a default rate limit of 1000 requests per minute per API key. This limit is configurable per customer on Enterprise plans, up to a hard ceiling of 50,000 requests per minute per key.

Requests exceeding the configured limit receive an HTTP 429 response with a Retry-After header. Burst traffic is smoothed using a token bucket algorithm with a bucket size equal to 10 percent of the per-minute limit.

Legacy note: the original 2021 beta release used a default rate limit of 200 requests per minute per API key before the limit was raised in the first general-availability release.

Authentication

Aurora supports OAuth2 client-credentials flow and static API keys for service-to-service calls. JWT bearer tokens signed with RS256 are accepted for user-delegated requests, with a maximum token lifetime of 1 hour before re-issuance is required.

Supported Regions

Aurora is currently deployed in two regions: US-East and EU-West. Cross-region failover is available on the Pro and Enterprise tiers, with an expected failover time of under 30 seconds.

Data residency guarantees mean that traffic routed through the EU-West region never transits US-East infrastructure, to support customers with regional compliance requirements.

Pricing

The Free tier includes 10,000 requests per month at no cost, with access to US-East only and no SLA guarantee. The Pro tier costs $99 per month and includes 1,000,000 requests, both regions, and a 99.9 percent uptime SLA.

The Enterprise tier is custom-priced and includes dedicated capacity, configurable rate limits, a 99.99 percent uptime SLA, and a dedicated support engineer.

Observability

Aurora exposes request-level logs, latency histograms, and error rates via a built-in dashboard, with a 30-day retention window on the Pro tier and a 1-year retention window on Enterprise.

Webhooks

Customers can register webhook endpoints to receive near-real-time notifications for rate-limit breaches, certificate expiry, and region failover events. Webhook delivery is retried up to 5 times with exponential backoff before being marked as failed.